

IR Assignment 4 Report

Step-wise Explanation of the code

1. Data Source Import Setup: This section of code sets up the environment for importing data from Kaggle. It defines constants and functions required to download and extract data from URLs provided in the `'DATA_SOURCE_MAPPING'` variable.

2. Data Download and Extraction: The script downloads the specified dataset from Kaggle using the URLs provided in `'DATA_SOURCE_MAPPING'`. It checks if the file is a ZIP or TAR file and extracts it accordingly.

3. Data Loading and Preprocessing: After downloading and extracting the dataset, the script loads it into a Pandas DataFrame. It then selects the first 20,000 rows for further processing.

4. Text Preprocessing: The script preprocesses the text data by performing the following steps:

- Lowercasing: Converts all text to lowercase to ensure consistency.
- Tokenization: Splits the text into individual words or tokens.
- Punctuation Removal: Removes punctuation marks from the text.
- Stopword Removal: Removes common stopwords like 'and', 'the', 'is', etc., using NLTK's list of English stopwords.
- Lemmatization: Reduces words to their base or root form using NLTK's WordNet lemmatizer. This step helps in standardizing the words (e.g., "running" becomes "run").

5. Stemming: The code also defines a function called `'lemmatize_text'` that lemmatizes tokens in the text. It then applies this function to the `'processed_text'` column of the DataFrame.

6. Duplicate Removal: Finally, the script removes duplicate entries based on the preprocessed text column to ensure unique samples in the dataset.

7. Imports: Import necessary libraries including pandas for data manipulation, numpy for numerical operations, nltk for natural language processing tasks, and re for regular expressions.

8. NLTK Resources: Downloads necessary NLTK resources such as word tokenizer, stopwords, and WordNet corpus.

9. Data Cleaning:

- **Drop NaN Rows:** Removes rows with NaN values in the 'Summary' column.
- **Lowercasing:** Converts the 'Summary' text to lowercase to ensure consistency.
- **Tokenization:** Splits the 'Summary' text into individual words or tokens using NLTK's word tokenizer.
- **Removing Punctuation:** Eliminates punctuation marks from the tokenized text.
- **Stopword Removal:** Filters out common stopwords from the text using NLTK's English stopwords list.

- Join Tokens: Joins the processed tokens back into text form.

10. Stemming and Lemmatization:

- Stemming: This section defines a function ``lemmatize_text`` to lemmatize tokens in the text using NLTK's WordNet lemmatizer. The lemmatized tokens are then joined back into text form.
- Apply Lemmatization: Applies the lemmatization function to the 'processed_summary' column of the DataFrame.

11. Remove Duplicates: Removes duplicate entries based on the processed summary text to ensure unique summarization samples.

12. Splitting into Train and Test Sets: Splits the data into training and testing sets using a 75-25 split ratio.

13. Creating DataFrames: Constructs pandas DataFrames for both the training and testing data, containing the processed text and summaries.

14. Reset Index: Optionally resets the index of the DataFrames for consistency.

15. Text Combination and Formatting:

- Combining Text: Combines the first 100 words of the processed text with the processed summary, separated by 'TEXT:' and 'SUMMARY:' respectively.
- Adding End Marker: Appends 'END' to the end of each combined text to indicate the end of the sequence.

16. Mounting Google Drive (Optional): If running on Google Colab, mounts Google Drive to save or load files.

17. Writing Samples to Text File:

- This section opens a text file in write mode and iterates through the samples in the 'processed_text' column of the 'train_data' DataFrame.
- Each sample is written to the file with a prefix "Sample {index}:" followed by the sample text.

18. Custom Dataset Class:

- Defines a custom dataset class ``CustomDataset`` inheriting from ``torch.utils.data.Dataset``.
- The class takes the file path, tokenizer, and block size as input parameters.
- In the constructor (``__init__``), it reads the text from the file, tokenizes it using the tokenizer, and creates examples of fixed block size from the tokenized text.
- Implements ``__len__`` and ``__getitem__`` methods required for a PyTorch dataset.

19. Load Dataset Function:

- Modifies the ``load_dataset`` function to use the custom dataset class (``CustomDataset``).
- It initializes an instance of ``CustomDataset`` with the provided parameters and returns it.

20. Load Data Collator Function:

- Defines a function ``load_data_collator`` to create a data collator for language modeling.
- It initializes a ``DataCollatorForLanguageModeling`` object with the provided tokenizer and whether MLM (masked language modeling) is used.

21. Training Function:

- Defines a function ``train`` to perform fine-tuning of the GPT-2 model on the custom dataset.
- It loads the dataset using the ``load_dataset`` function and the data collator using the ``load_data_collator`` function.
- Initializes the GPT-2 model and sets up training arguments including output directory, batch size, number of epochs, etc.
- Instantiates a ``Trainer`` object with the model, training arguments, data collator, and train dataset.
- Calls the ``train`` method of the trainer to start the training process.
- Saves the trained model.

22. Training Parameters Setup:

- Defines parameters such as the path to the training file, model name, output directory, batch size, number of epochs, and save steps.

23. Training Invocation:

- Invokes the ``train`` function with the provided parameters to start the training process.

24. Function Definitions:

- ``load_model(model_path)``: Loads the fine-tuned GPT-2 model from the specified ``model_path``.
- ``load_tokenizer(tokenizer_path)``: Loads the GPT-2 tokenizer from the specified ``tokenizer_path``.
- ``generate_text(sequence, max_length)``: Generates text (summary in this case) given a starting ``sequence`` and a maximum length (``max_length``) using the loaded model and tokenizer. It utilizes the ``generate`` method of the GPT-2 model with specific parameters such as sampling, maximum length, and top-k sampling.

25. Text Generation Loop:

- Initializes an empty list ``predicted_summaries`` to store the generated summaries.
- Initializes an index variable ``i`` to iterate through the ``processed_text`` column of the ``test_data`` DataFrame.
- Enters a while loop that continues as long as the index ``i`` is less than the length of the ``processed_text`` column.
- Retrieves a ``sequence`` from the ``processed_text`` column at index ``i``.
- Calculates the ``max_len`` for the sequence by adding 100 to the number of words in the sequence.
- Calls the ``generate_text`` function to generate a summary for the sequence with the specified ``max_len``.

- Extracts the generated summary from the output and appends it to the `predicted_summaries` list.
- Increments the index `i` to move to the next sequence in the loop.

26. Explanation:

- The code essentially generates summaries for each sequence in the `test_data` DataFrame by iterating through its `processed_text` column.
- For each sequence, it generates a summary using the fine-tuned GPT-2 model, considering a maximum length of the sequence plus 100 additional words.
- The generated summary is then appended to the `predicted_summaries` list.
- Finally, the loop continues until summaries are generated for all sequences in the `test_data` DataFrame.

Results

Rogue scores

```
{'rouge-1': {'r': 0.7946693121693122, 'p': 0.818880952380952, 'f': 0.8007613221107203}, 'rouge-2': {'r': 0.6603968253968252, 'p': 0.6641111111111112, 'f': 0.656248801718073}, 'rouge-l': {'r': 0.7946693121693122, 'p': 0.818880952380952, 'f': 0.8007613221107203}}
```